

CSE 231 Project 1

Michael Barrow
mbarrow@eng.ucsd.edu
University of California, San Diego

Jules Testard
Second Institution

Costas Zafarias
aweff

Abstract

This report describes the implementation and experimental results of custom LLVM passes applied to supplied benchmark programs. The high level concept of each pass is discussed, followed by key implementation details. Finally we offer conjecture on experimental application of these passes.

1 Introduction

LLVM is a collection of modular compiler tools. The project's original goal was a flexible compilation strategy for arbitrary programming languages able to perform both static and dynamic compilation.

This language flexibility is mainly achieved with a common intermediate program representation during the compilation known as LLVM byte code. Source language and target machine independent optimisations can be made to the LLVM byte code, which is the 231 course project subject of interest.

Our prescribed source language is C++ and our target architecture is x86. The project goal is to profile a set of LLVM compiled C++ benchmarks programs using the **pass** feature of the **opt** LLVM module. Three pass functionalities are required:

- Collecting static instruction counts
- Collecting dynamic instruction counts
- Profiling branch bias

The required depth of understanding and proficiency in LLVM module code increases incrementally with each of these three functionalities. Therefore the report is split into three mini-reports to focus on the self contained lessons learned

from implementing them. Mini-reports have the following sub sections:

- **Instrument description:** A high level description of an *algorithm* used to implement functionality
- **Instrument implementation summary:** Details on key LLVM concepts and api's used to implement the functionality *algorithm* as an *opt pass*
- **Benchmark analysis:** Conjecture based running the *opt pass*

2 Static instruction count

Problem statement: Write a pass that counts the number of static instructions in a program. **Problem instance:**

- output total number of instructions
- output per-instruction count

2.1 Pass description

Because the functionality is a static analysis, it is sufficient to run a pass on the compiled code of each benchmark. We store instruction op codes and their corresponding count in a C++ map structure. Each time an instruction is found in the source code, it is either added to an existing map entry or a new entry is created and initialized to 1.

A high level algorithm description would be:

```

Input:  $M, I$ 
1  $t \leftarrow 0$ 
2 forall the  $i \in I$  do
3   if  $M.containsKey(i)$  then
4      $M.valueForKey(i) += 1$ 
5   else
6      $M.insertKeyValuePair(< i, 1 >)$ 
7 forall the  $keyValuePair \in M$  do
8    $print("Found "keyValuePair.value() "counts of:$ 
9      $keyValuePair.key())$ 
10   $t += keyValuePair.value()$ 
11  $print("total instructions: "t)$ 

```

Algorithm 1: Static instruction count algorithm

where:

I is an input program instruction list in LLVM byte code (.bc) format

i is an individual instruction within I

M is a C++ map of the form `jsstring,int;`

2.2 Static pass implementation summary

Besides a working proficiency in C++, translating the algorithm to LLVM required an understanding of:

- how to build and run an opt module on a target benchmark program
- how to output data in human readable format from an opt module
- how instructions are represented by LLVM
- how instructions are accessed by opt modules

The most challenging of the above was understanding how instructions are represented by LLVM and accessing them from opt modules.

Our opt pass is able to iterate through the source module by implementing the `runOnModule()` virtual function of a `ModulePass` class.

```
virtual bool runOnModule(Module &M)()
```

Accessing instructions was accomplished by iterating through the **input source Module** using an **inst_iterator**. The Module contains all instructions from the compiled benchmark and the `inst_iterator` points at individual instructions.

```

for (Module::iterator m = M.begin(), e = M.end())
  for (inst_iterator I = inst_iterator::begin(), E = inst_iterator::end())

```

2.3 Benchmark analysis

3 Dynamic instruction count

Problem statement: Write a pass that instruments an input benchmark to count the number of times each instruction executes **Problem instance:**

- count program instructions at runtime
- output-per instruction count

3.1 Pass description

The pass differs from the prior pass in that it instruments the target byte code to perform an analysis on itself. After instrumentation, the benchmark program executes and dynamically counts instructions, outputting the total at termination. This dynamic analysis is therefore two phase with phase 1 being the opt pass and phase 2 being program execution.

First we describe the intended instrumented binary diagrammatically in Figure1




Figure 1: Required bytecode transformation

With a target binary defined we describe a code transformation that the opt pass must perform:

where:

I is an input program instruction list in LLVM byte code (.bc) format

i is an individual instruction iterator I

The `count()` function can be recycled from section the *static count* pass. Similarly the `Print()` function can be recycled.

Input: I

```
1  $x \leftarrow 0$ 
2  $i \leftarrow \text{ref}(I.\text{at}(0))$ 
3 while  $i + 1 \neq \text{do}$ 
4    $i' \leftarrow \text{Count}(\text{valAt}(i))$ 
5    $i.\text{insertBefore}(i')$ 
6    $i += 1$ 
7  $i \leftarrow \text{ref}(I.\text{last}())$ 
8  $i.\text{insertBefore}(\text{print}())$ 
```

Algorithm 2: Dynamic instruction count instrumentation pass

3.2 Instrumentation pass implementation summary

Besides the skills used for the static pass, implementation of the dynamic pass presented challenges of inserting instrumentation functions to the test bench code and further augmenting that code to call inserted functions.

Ensuring instrumentation functions were in scope during the opt pass was documented in the one of the project brief "hints" sections. Counting sorted lists of instructions was a problem solved during the static instruction count problem. Therefore the major challenges were:

- Finding a function signature in LLVM byte code
- Inserting a call to a function signature above a byte code instruction
- Passing parameters to an inserted function call

Finding a function signature **COSTAS IMPROVED THIS, I WILL SAY NO MORE**

Inserting a function call

4 Profiling branch bias

4.1 Conclusions

The project was completed successfully. Learning outcomes were a working understanding of LLVM optimiser passes. We are capable of transforming and instrumenting code. We have developed an efficient method to develop, test and debug the LLVM source tree with a debug api and symbol indexed code base via the eclipse CDT. Understanding how to instrument code has provided a powerful tool for analysis of any optimisation heuristics applied

by other opt passes. Instrumentation will allow us to both formulate and test heuristics based on the performance of compiled code.

To conclude, this project has provided us with a solid foundation for the second upcoming project for CSE 231

References

Notes

¹Remember to use endnotes, not footnotes!